

# Speech Signal Processing and Speech Enhancement Summer 2023

Denoising Diffusion on Speaker Embeddings


Akshat Gupta

# Introduction


- Generative Deep Learning Model
- Eg. Audio Generation, Image Generation
- Uses Markov Chain to learn the distribution of data- mean, variance



Diffusion



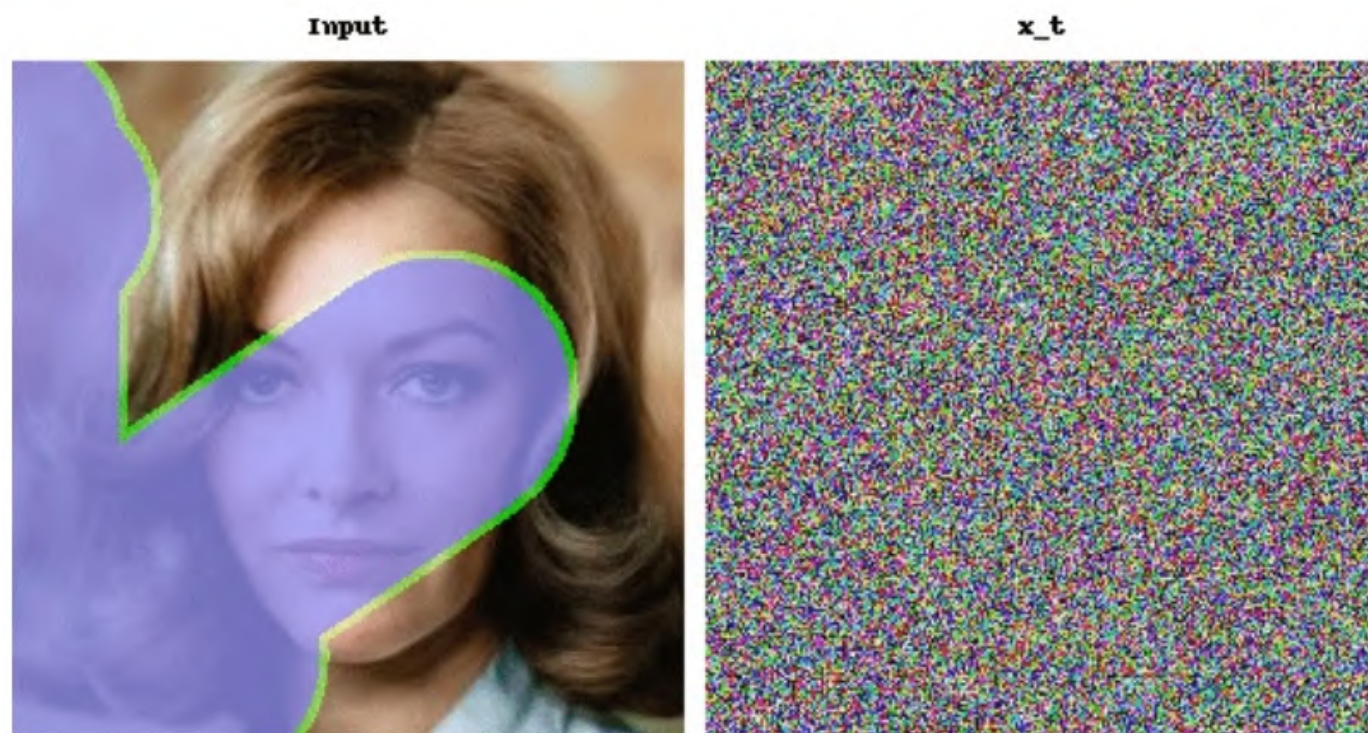
# Motivation





# Motivation

- Image Inpainting- Repaint
- Anomaly Detection - AnoDDPM



RePaint

# Motivation

- Natural Language Processing- Text Synthesis
- Multi-Modal Learning- Text to Image - Imagen, DALL-E, Text to Audio, Diff-TTS

A group of teady bears in suit in corporate office celebrating the birthday





# Diffusion Model



# Popularity of Diffusion Model



Jonathan Ho

Unknown affiliation  
Verified email at berkeley.edu - [Homepage](#)  
[Artificial Intelligence](#) [Machine Learning](#)

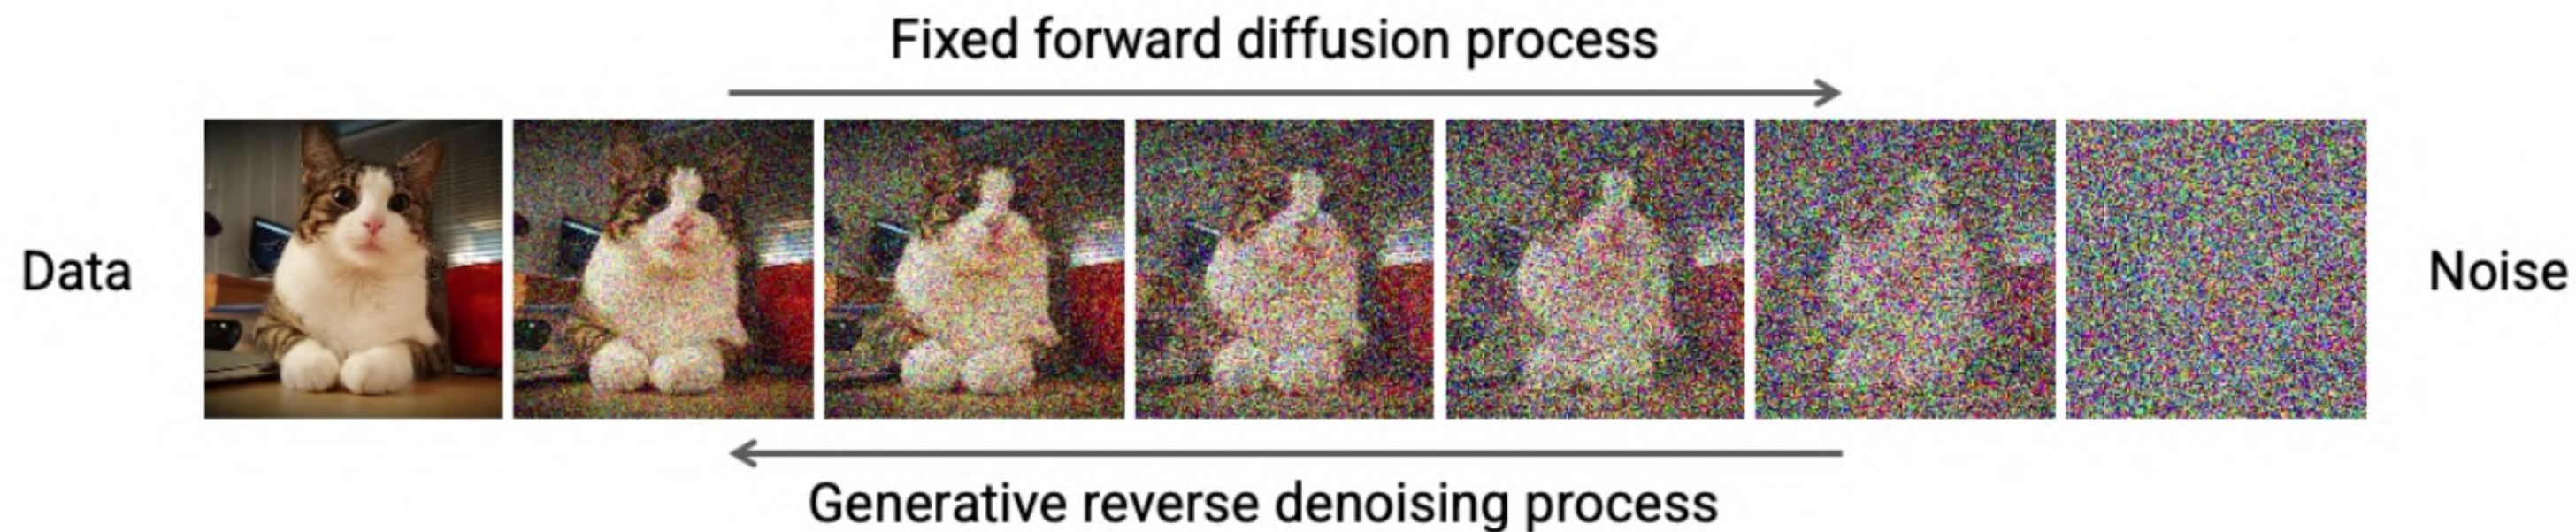


TITLE	CITED BY	YEAR
<a href="#">Generative adversarial imitation learning</a> J Ho, S Ermon Advances in Neural Information Processing Systems, 4565-4573	2273	2016
<a href="#">Evolution strategies as a scalable alternative to reinforcement learning</a> T Salimans, J Ho, X Chen, S Sidor, I Sutskever arXiv preprint arXiv:1703.03864	1314	2017
<a href="#">Denoising diffusion probabilistic models</a> J Ho, A Jain, P Abbeel Advances in Neural Information Processing Systems 33, 6840-6851	1032	2020



# How DDPM Works?

- Learning to generate using denoising
- Two processes:
  - Forward Diffusion Process: adds noise to image
  - Reverse Denoising Process: generate new data





# Forward Diffusion Process

- Forward Process in T steps

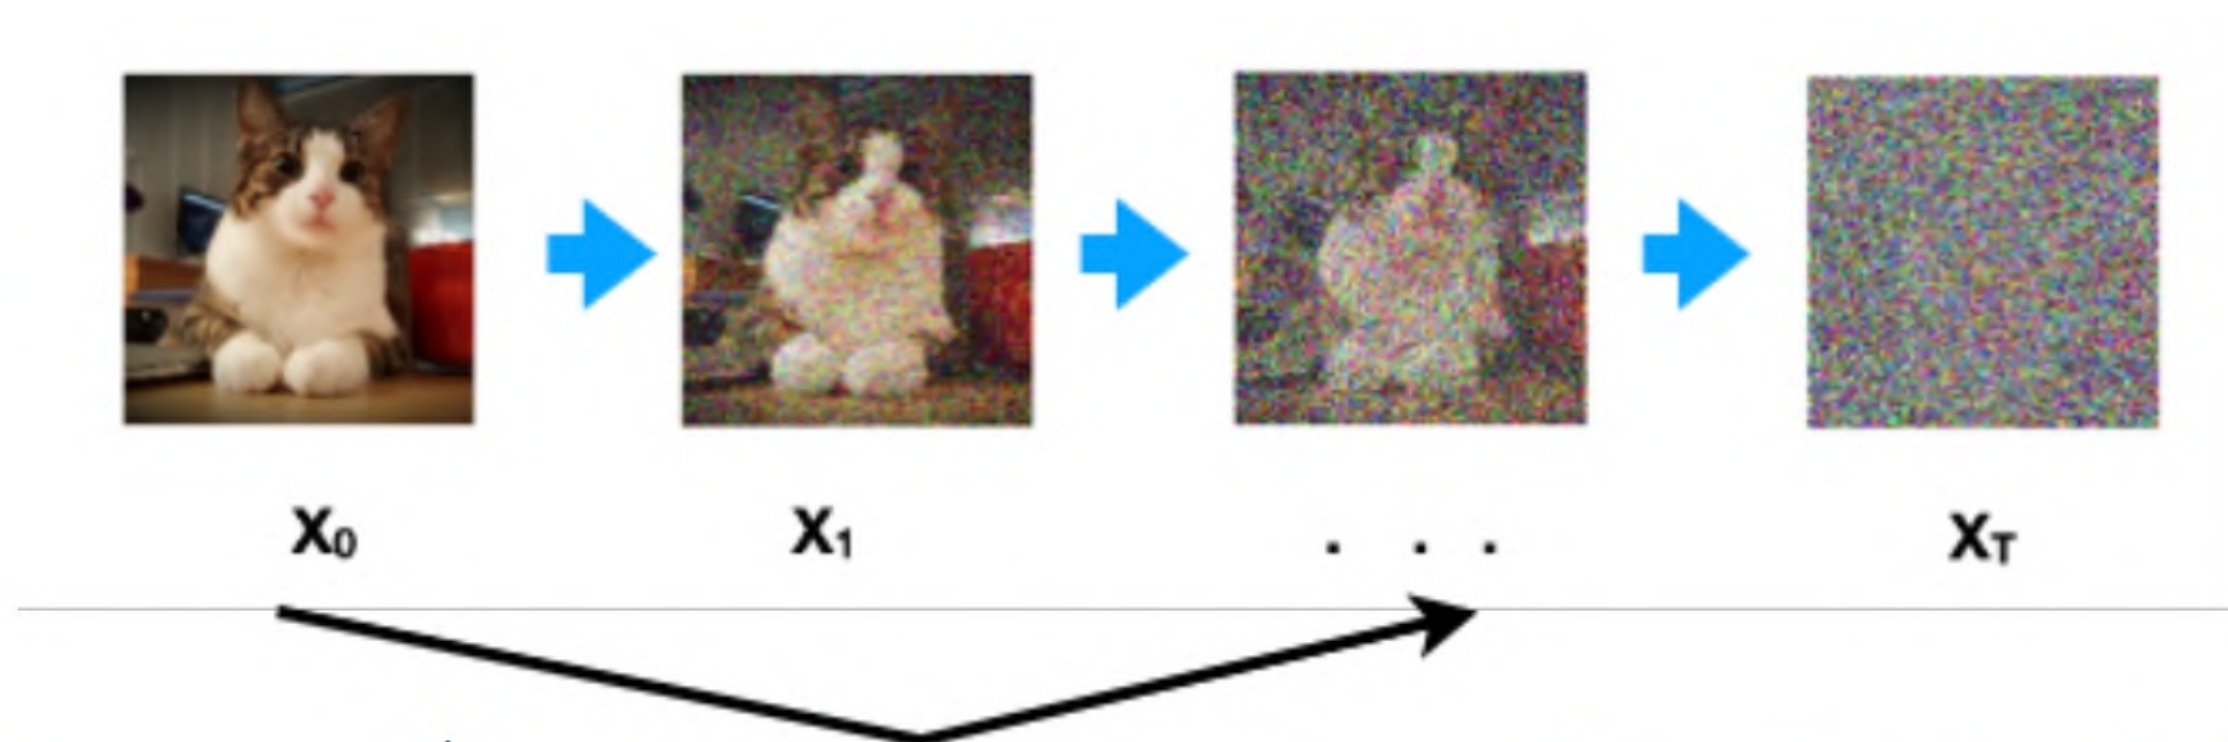


Source: <https://cvpr2022-tutorial-diffusion-models.github.io/>

$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \beta_t \mathbf{I}) \longrightarrow q(\mathbf{x}_{1:T} | \mathbf{x}_0) = \prod_{t=1}^T q(\mathbf{x}_t | \mathbf{x}_{t-1}) \quad (1)$$

- $\beta_t$  : variance at timestamp  $t$ ,  $\sqrt{1 - \beta_t}$  : mean at timestamp  $t$
- $q(x_t | x_{t-1})$  : distribution at each timestamp  $t$ ,  $q(x_{1:T} | x_0)$  : joint distribution

# Diffusion Kernel



$$\text{Define } \bar{\alpha}_t = \prod_{s=1}^t (1 - \beta_s) \longrightarrow q(\mathbf{x}_t | \mathbf{x}_0) = \mathcal{N} \left( \mathbf{x}_t; \sqrt{\bar{\alpha}_t} \mathbf{x}_0, (1 - \bar{\alpha}_t) \mathbf{I} \right) \quad (2)$$


For sampling (reparameterization trick)

$$\mathbf{x}_t = \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{(1 - \bar{\alpha}_t)} \varepsilon \text{ where } \varepsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I}) \quad (3)$$

$\beta_t$  values scheduled (variance or noise schedule)


$$\bar{\alpha}_T \rightarrow 0 \text{ and } q(\mathbf{x}_T | \mathbf{x}_0) \approx \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I}) \quad (4)$$

# How Diffusion happens?

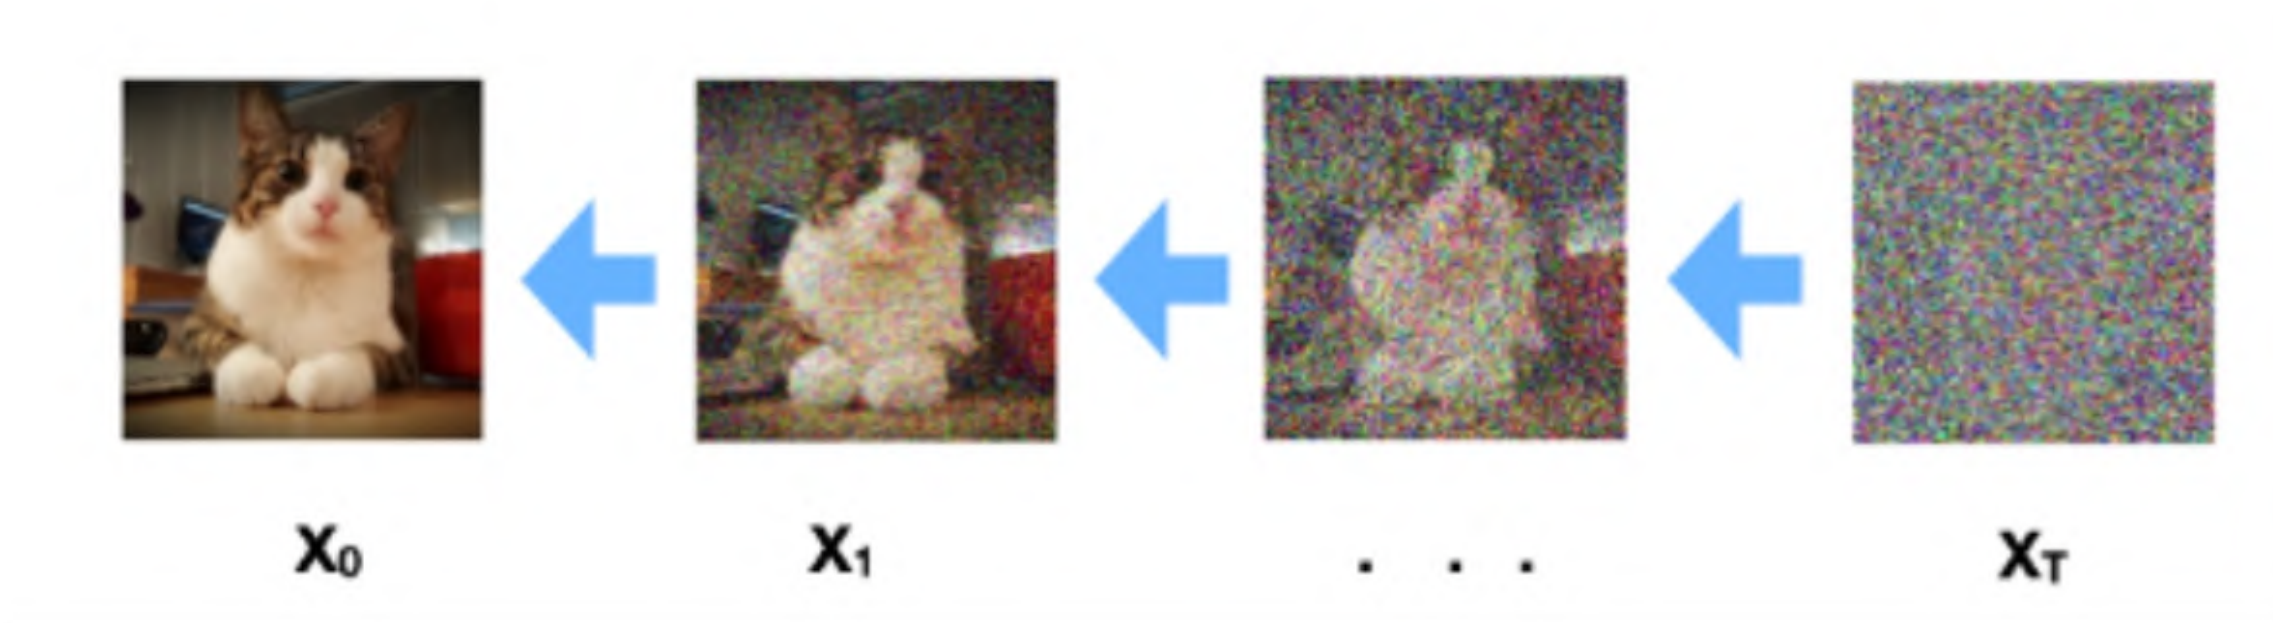




# How Diffusion happens?



# Reverse Diffusion Process




$$p(\mathbf{x}_T) = \mathcal{N}(\mathbf{x}_T; \mathbf{0}, \mathbf{I})$$

$$p_{\theta}(\mathbf{x}_{t-1} | \mathbf{x}_t) = \mathcal{N}(\mathbf{x}_{t-1}; \mu_{\theta}(\mathbf{x}_t, t), \sigma_t^2 \mathbf{I})$$


Trainable Network  
(U-net, Denoising Autoencoder)

- $p_{\theta}(x_{t-1}|x_t)$  : denoising diffusion probability

# What happens to distribution?



Forward (Diffusion) Process




Reverse (Denoising) Process

Source: <https://cvpr2022-tutorial-diffusion-models.github.io/>



# Network architecture

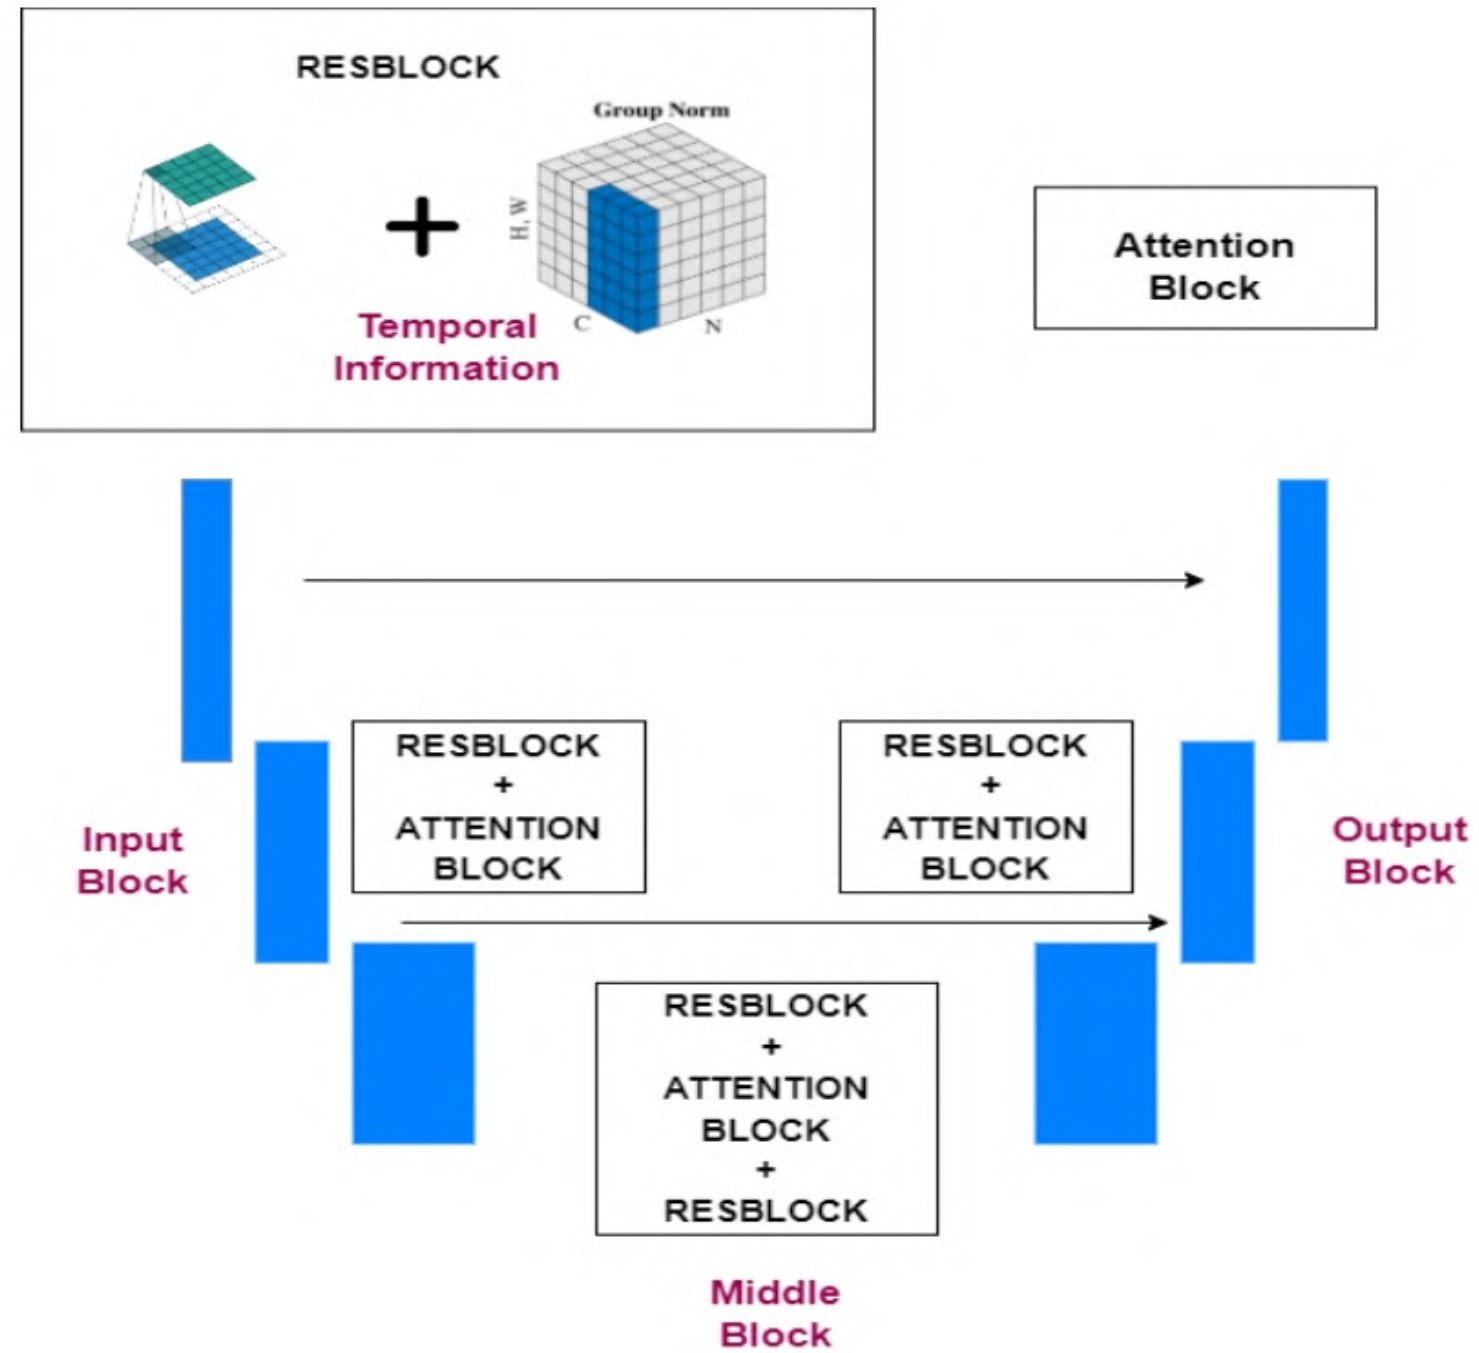
- U-NET architecture
- Train model predicts  $\epsilon_\theta$  (noise)



$$L_{\text{simple}} := \left\| \epsilon - \epsilon_\theta \left( \sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t \right) \right\|^2$$


Source: <https://cvpr2022-tutorial-diffusion-models.github.io/>

# U-NET architecture



# Evaluation of DDPM

- Sample qualitative metrics
  - **Fidelity** - how realistic generated image is
  - **Diversity** - fake samples capture the variation
- Sample Quantitative metrics
  - **Inception Score (IS)**: sharpness and diversity, synthetic image  $\neq$  real images
  - **Frechet Inception Distance (FID)**: Distribution of synthetic images and real images
- **Mean and Variance** of the learned distributions and compare it to the real distribution





# Introduction to SpeakEMB

- Denoising Diffusion Probabilistic Models (DDPM) have emerged as prominent generative models.
- These models have been implemented predominantly on image synthesis
- Gradually moving to the domain of speech
- In this project, I aspire to use DDPM to generate Speaker Embedding.


# DDPM on Speaker Embeddings

- Denoising Diffusion Model uses diffusion to add Gaussian Noise at every time step to the Speaker Embedding
- Denoising to remove the Gaussian Noise and generate a new set of Speaker Embedding

Forward Process: Diffusion

$$q(x_0) \rightarrow q(x_T) = \mathcal{N}(x_T, 0, I)$$
$$q(x_t|x_{t-1}) = \mathcal{N}(x_t, x_{t-1}\sqrt{1 - \beta_t}, I\beta_t)$$

The clean speaker embedding is converted into an isotropic Gaussian distribution in a step-by-step diffusion process.




# DDPM on Speaker Embeddings

- When denoising is done on an isotropic Gaussian distribution, a new embedding can be generated

Backward Process: Denoising

$$p(x_0) = q(x_0) \leftarrow p(x_T) = \mathcal{N}(x_T; 0; I)$$
$$p(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; f_\mu(x_t, t); f_\Sigma(x_t, t))$$

In the reverse process, it tries to gradually restore the clean input by predicting and removing the noise introduced in each step of the diffusion process.




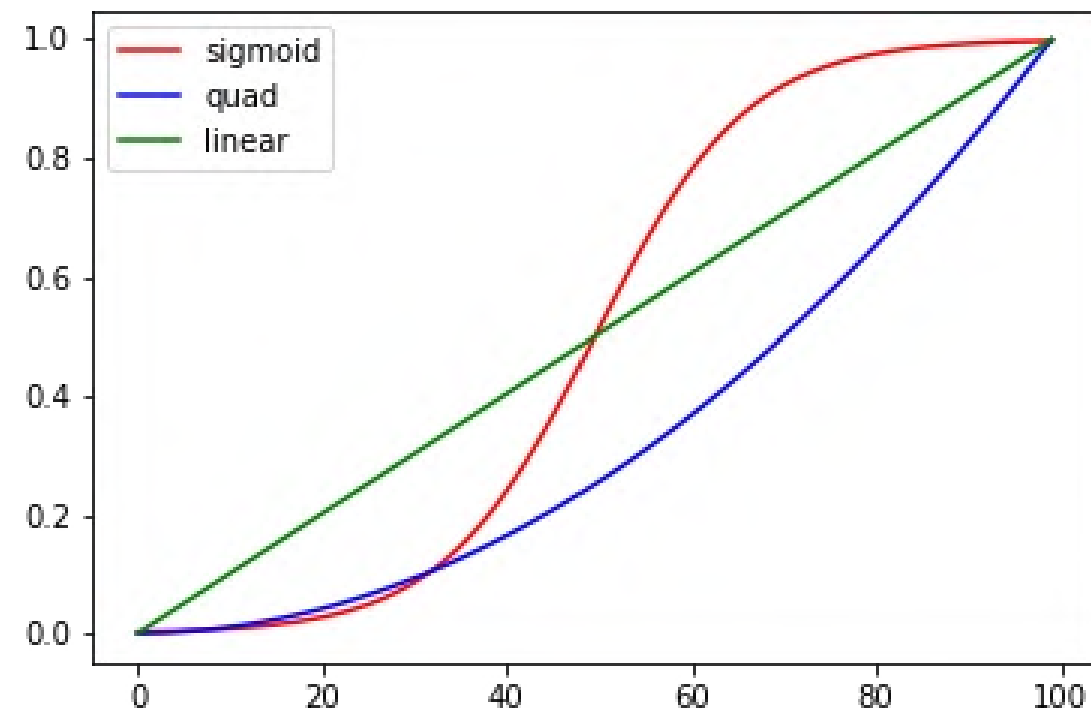


# Methods

- The baseline model is implemented on Denoising Diffusion Probabilistic Model to Generate Speaker Embedding
- Integrating the model architecture to a three layer fully connected Neural Network. Mount L1 loss in the baseline model and the following section.three-layer

# Variance Scheduling

- Four different Variance Schedulers of Linear, Sigmoid and Quadratic bearing ranges from  $e-5$  to  $e-2$
- Cosine bearing ranges from 0 to 1 on our Baseline Model with L1 loss
- The improvements in Cosine Variance Schedulers have shown promising results by Nichol et al.



# Dataset

- Extracted speaker embedding on LibriSpeech dataset with Speechbrain Framework with ECAPA-TDNN Model, Hugging Face
- The extracted speaker embedding has different dimensionalities of 64, 128 and 704

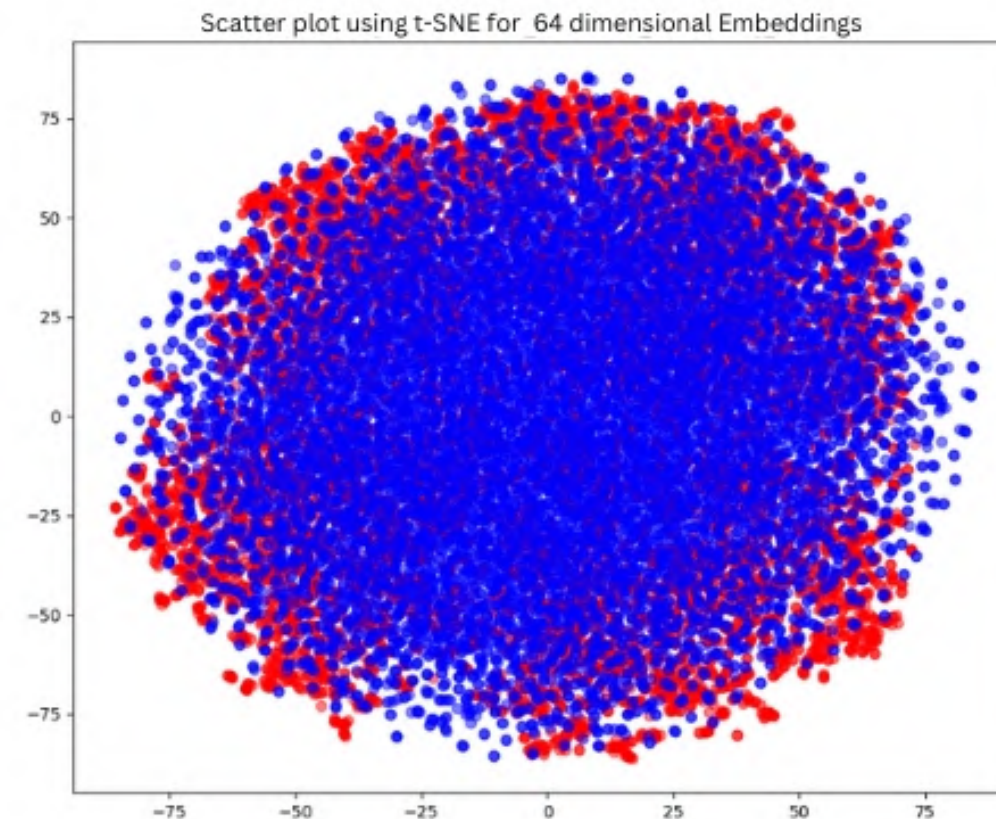


# Experiments

- The experiments are performed on three different sets of embedding having dimensions 64, 128 and 704
- The TSNE plot consists of the red data points which are corresponding to the original set of embedding and the blue data points which are corresponding to the generated set of embedding.
- The loss graph indicates learning and the convergence of the model.
- To visualize, we perform dimensionality reduction on the embedding and monitor the formation of clusters.


# 64 Dimension Experiment

- The fundamental experiment was performed on 64 dimension embedding until the generated embedding could
- be mapped into some clusters.
- The formation of these clusters is an illustrative explanation of how well the generation of the speaker embedding has been executed.



# 128 Dimension Experiment

- With the increase in the number of the dimension of the speaker embedding from 64 to 128, the cluster formation of the generated embedding has been prominent.
- The blue clusters are constrained in the centre which shows that there has been a significant amount of learning but it definitely requires a large training cycle to fairly replicate the distribution of the original embedding.

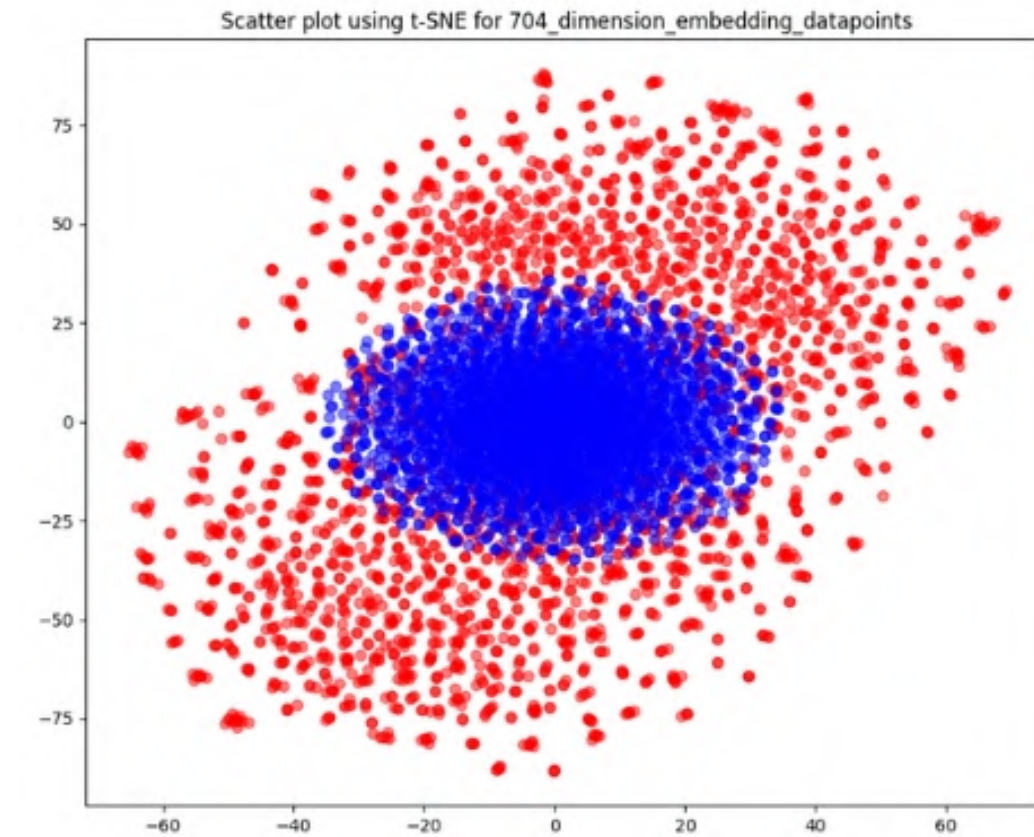




# 704 Dimension Experiment

- The 704 generated embedding are lying in the centre of the latent space.
- This is simply because the distribution of the generated embedding still has not completely learned the distribution of the original embedding.
- The reason behind this is the increase in the dimensionality of the embedding, as it would require an increase in the number of training steps of the model.
- The generated 704 embeddings can be used to evaluate our performance on the real-world application as we used the embedding of dimension 704 and tested the performance on a text-to-speech IMS Toucan framework by the University of Stuttgart.
- Observed that every generated embedding is distinct.

# 704 Dimension Experiment



# Results

- The difference in Mean and Variance between Original and Generated Embedding help in evaluating how well, the generated embedding is learnt.
- Ran all sets of experiments only for 500 steps to compare the convergence on a different set of embedding.

Dimensions	Data points	Steps	Mean Difference	Variance Difference
64	500	500	$2.38e - 07$	-4.4288
64	1000	500	$1.31e - 06$	-10.16
128	500	500	$1.19e - 07$	0.799
128	1000	500	$-1.86e - 08$	0.389
704	500	500	$1.79e - 07$	-0.7833
704	1000	500	$7.22e - 16$	-1.4066

**Table 1:** Difference in Mean and Variance between Original and Generated Embedding.



# Conclusion

- DDPM are successful at generating unseen speaker embedding
- Perform Controllability of Unseen Speaker Embedding with Eigenvector rotation which will lead to the generation of Speaker embedding bearing specific age, gender, etc

# References

- [1] Brecht Desplanques, Jenthe Thienpondt, and Kris Demuynck. ECAPA-TDNN: emphasized channel attention, propagation and aggregation in TDNN based speaker verification. In Helen Meng, Bo Xu, and Thomas Fang Zheng, editors, Interspeech 2020, pages 3830–3834. ISCA, 2020.
- [2] Jonathan Ho, Ajay Jain, and Pieter Abbeel. Denoising diffusion probabilistic models. *Advances in Neural Information Processing Systems*, 33:6840–6851, 2020.
- [3] Florian Lux, Julia Koch, Antje Schweitzer, and Ngoc Thang Vu. The IMS Toucan system for the Blizzard Challenge 2021. In *Proc. Blizzard Challenge Workshop*, volume 2021. Speech Synthesis SIG, 2021.
- [4] Alexander Quinn Nichol and Prafulla Dhariwal. Improved denoising diffusion probabilistic models. In *International Conference on Machine Learning*, pages 8162–8171. PMLR, 2021.

Thank you  
Questions?